

Hybrid Mechanistic, Bayesian Inference & Optimization System for Passion-Fruit Pectin Extraction

Research & Experimental Design Blueprint

Contents

1. Executive Summary	3
2. Generalized Methodological Framework	3
2.1. The Three Pillars of System Intelligence	3
2.2. The Nine Core Components	4
2.3. The Core Abstraction	4
3. Mechanistic Reactor Model (V1.3)	5
3.1. Compartmental Cascade	5
3.2. Kinetic Equations (Mass & Moments)	5
3.3. Parameter Ontology	6
4. Structural Identifiability & OED	6
4.1. Bayesian Optimal Experimental Design (OED)	6
5. Hierarchical Measurement Architecture	7
5.1. The Preparation Model & Precipitation Bias	7
5.2. The Proxy Observation Model	8
5.3. Proxy Calibration & Acceptance Criteria	8
5.4. Bayesian Integration	8
6. Physical Implementation & Minimum Viable Laboratory (MVP)	9
6.1. The Lean V1 Sensor Stack	9
6.1.1. Continuous Sensor Priority Justification	9
6.1.2. The Empirical Sensor Fusion Philosophy	10
6.2. Equipment Sourcing Strategy	10
6.3. Analytical Discipline & Traceability	11
7. Conclusion & Execution Map	11
8. Appendix A: Primer on Bayesian Inference	12
8.1. A.1 The Conceptual Shift: From Point Estimates to Probability	12
8.2. A.2 Practical Application in the Pectin Digital Twin	12
8.2.1. 1. Quantifying Uncertainty for Process Optimization	12

8.2.2. 2. The Hierarchical Measurement Model	13
8.3. A.3 Computational Execution	13

1. Executive Summary

This document synthesizes the complete research architecture for building a digital twin of passion-fruit pectin extraction. The architecture bridges a strict mechanistic kinetic model with a Bayesian Optimal Experimental Design (OED) framework. To solve the fundamental bottleneck of expensive laboratory analytics, the system employs a hierarchical proxy measurement model, allowing dense time-series data collection using cost-effective chemometric and rheological sensors.

To provide a clear distinction between the methodology’s conceptual sequence and its current state of progress, this blueprint is structurally divided. After outlining the complete 9-component methodology, **Part I** explicitly reports on the steps that have already been executed via modeling and synthetic data simulations (Components 1-3). **Part II** outlines the pending physical laboratory protocols required to execute the remaining steps in reality (Components 4-5).

The sequence of this research project is explicitly defined as:

Mechanistic Model → Synthetic Recovery → OED → Proxy Calibration
→ Physical Experiment → Bayesian Update → Optimization

2. Generalized Methodological Framework

While this document details the specific application of pectin acid hydrolysis, the underlying architecture represents a highly generalizable **Hybrid Mechanistic–Statistical Experimental Learning System**. This methodology applies broadly to bioprocessing, materials synthesis, and reaction engineering where experiments are expensive and optimization is critical.

2.1. The Three Pillars of System Intelligence

The framework explicitly separates three distinct computational responsibilities:

Mechanistic Model → What can happen? (Physics/Chemistry)
Bayesian Inference → What do we currently believe? (Uncertainty)
OED / Active Learning → What should we measure next? (Information)

Machine learning surrogates serve strictly to make these three operations computationally affordable.

2.2. The Nine Core Components

The complete conceptual architecture consists of nine integrated components:

1. **Mechanistic Process Model:** Represents the underlying physical system ($x_{\{t+1\}} = f(x_t, u_t, \theta)$), providing interpretability, extrapolation structure, conservation laws, and a place to encode prior scientific knowledge.
2. **Identifiability Analysis:** Determines what can actually be learned from available measurements using sensitivity matrices ($S = \frac{\partial y}{\partial \theta}$).
3. **Optimized Experimental Design (OED):** Selects experiments that maximize information gain ($x^* = \arg \max_x \mathcal{I}(x)$) using formal optimality criteria (e.g., D-optimality).
4. **Hierarchical Proxy Measurement Model:** Models the entire observation chain (True State $\rightarrow$ Preparation $\rightarrow$ Proxy $\rightarrow$ Observation) rather than assuming perfect sensors, turning expensive assays into calibration anchors for dense, cheap measurements.
5. **Bayesian Parameter Inference:** Infers the full posterior probability ($p(\theta \mid D)$) rather than point estimates, quantifying uncertainty in both parameters and future predictions ($p(y^* \mid x^*, D)$).
6. **Surrogate ML Modeling:** Learns the expensive forward mapping ($\hat{f}(x, \theta) \approx f(x, \theta)$) to accelerate inference and optimization.
7. **Optimization Under Uncertainty:** Replaces standard expected-value maximization with robust decision-making (e.g., maximizing the probability of an acceptable outcome subject to quality/reliability constraints).
8. **Sequential Active Learning:** Closes the loop. The system does not run once; it continuously iterates: Model $\rightarrow$ Optimal Experiment $\rightarrow$ Data $\rightarrow$ Bayesian Update $\rightarrow$ Next Experiment.
9. **Model Validation & Discrepancy:** Acknowledges that $y = f(x, \theta) + \delta(x) + \varepsilon$, formally modeling the systematic structural mismatch ($\delta(x)$) between the mathematical equations and physical reality to enable true model criticism.

2.3. The Core Abstraction

The deepest abstraction of this system is not “use ML to optimize a reactor.” It is:

Build a mechanistic model, quantify what you don’t know, design the cheapest experiments that maximally reduce that uncertainty, use inexpensive measurements wherever possible through calibrated observation models, and close the loop by updating the model and selecting the next experiment.


PART I: EXECUTED STEPS (COMPUTATIONAL PHASE)

Reporting on the completion of Components 1–3 via numerical simulations and synthetic data

3. Mechanistic Reactor Model (V1.3)

3.1. Compartmental Cascade

The extraction is modeled as a 4-pool cascading mass transfer system. The cascade structure forces a critical structural tradeoff between extraction speed and product degradation.



3.2. Kinetic Equations (Mass & Moments)

The state equations combine strict normalized mass conservation $[\text{kg}/\text{kg}_{\text{dry,feed}}]$ with mass-weighted moments to track intensive properties (M_w, DE) dynamically.

$$\frac{dP_{\text{sol}}}{dt} = k_{\text{ext}} P_{\text{matrix}} - k_{\text{hyd}} P_{\text{sol}}$$

$$\frac{dP_{\text{lowMW}}}{dt} = k_{\text{hyd}} P_{\text{sol}} - k_{\text{deg}} P_{\text{lowMW}}$$

Kinetics are parameterized via an Arrhenius reference-state formulation ($T_{\{\text{ref}\}} = 353.15 \text{ K}$, $pH_{\{\text{ref}\}} = 2.0$) to dramatically improve the numerical conditioning of the Bayesian inference.

3.3. Parameter Ontology

To avoid mathematical ambiguity, all parameters are strictly assigned an ontology classification.

Symbol	Meaning	Units	V1 Range	Classification
T	Extraction Temperature	$^{\circ}\text{C}$	$45 - 100 \text{ }^{\circ}\text{C}$	Controlled Process Variable
pH	Kinetic Acidity	pH	$1.2 - 4.5$	Controlled Process Variable
t	Extraction Time	min	$5 - 120 \text{ min}$	Controlled Process Variable
$d_{\{50\}}$	Particle Size	μm	$149 - 600 \mu\text{m}$	Measured State
$M_{\{w,\text{matrix}\}}$	MW of matrix pectin	Da	6.54×10^5	Informative Prior
$k_{\{i,\text{ref}\}}$	Reference Rates	min^{-1}	TBD	Calibrated
$E_{\{i\}}$	Activation Energies	J/mol	TBD	Calibrated
$n_{\{i\}}$	Acidity Orders	-	1.0	Fixed V1
α	Particle Size Scaling	-	$0.5 - 2.0$	Calibrated

4. Structural Identifiability & OED

The serial cascade induces a severe identifiability problem: a fast extraction/fast degradation trajectory looks mathematically identical to a slow extraction/slow degradation trajectory if observed only at the process endpoint.

4.1. Bayesian Optimal Experimental Design (OED)

By computing the Fisher Information Matrix (FIM) over the 9-dimensional parameter space using synthetic data simulations, we designed an exact 10-run protocol that minimizes the determinant of the posterior covariance matrix (D-optimal design).

The OED analysis strongly favors time-resolved observation, particularly at early and intermediate residence times, to break the $k_{\{\text{hyd}\}} \leftrightarrow k_{\{\text{deg}\}}$ correlation. It pushes the boundary conditions of the operating space (e.g., $95 \text{ }^{\circ}\text{C}/\text{pH } 1.5$ vs $50 \text{ }^{\circ}\text{C}/\text{pH } 3.0$). For the final 10-run protocol utilizing 4-point time-resolved tracking, the selected protocol substantially improves the conditioning of the inverse problem, eliminating near-null sensitivity directions and achieving a condition number of $\kappa \approx 27.0$.

The explicit 10-run protocol defined by the solver is as follows (distributed across 2 baseline feedstock batches):

Run	T (°C)	pH	$d_{\{50\}}$ (μm)	Sample Times (min)
1	95	1.5	600	15, 45, 90, 120
2	50	1.5	150	15, 45, 90, 120
3	95	2.5	150	15, 45, 90, 120
4	50	1.5	600	15, 45, 90, 120
5	95	2.0	600	15, 45, 90, 120
6	50	2.0	150	15, 45, 90, 120
7	95	2.5	600	15, 45, 90, 120
8	50	1.5	300	15, 45, 90, 120
9	95	3.0	150	15, 45, 90, 120
10	95	2.0	300	15, 45, 90, 120

PART II: PENDING STEPS (PHYSICAL EXECUTION)

Proposed laboratory architecture to execute Components 4–5

5. Hierarchical Measurement Architecture

The time-resolved OED campaign requires 6 samples per run. Performing conventional SEC-MALLS (for M_w), Titration (for DE), and mHDP Colorimetry (for GalA purity) on every aliquot is economically unviable.

We solve this using a Bayesian Proxy Measurement Model. The physical experiment is segmented into a preparation layer and an observation layer.

5.1. The Preparation Model & Precipitation Bias

Direct measurement of raw reactor liquor is strongly confounded by the extraction matrix (unreacted acid, trace metal ions, simple sugars) and is therefore not the preferred basis for routine proxy measurements. All proxy samples undergo **Standardized Alcohol Precipitation**.

This is not a neutral step. Ethanol fractionation can selectively recover high- M_w pectin while leaving low- M_w fragments in solution. We model this recovery bias as $\eta_{\{\text{precipitation}\}} =$

$f(M_w, DE)$. The proxy must be calibrated on the exact same precipitated material to absorb this physical bias natively.

5.2. The Proxy Observation Model

We deploy two ultra-fast, low-cost sensors on the standardized precipitate:

1. **FTIR/DRIFTS Spectroscopy:** A chemometric PLS model maps the full mid-IR spectrum to DE and GalA fraction.
2. **Capillary Viscometry:** The flow time (t_{flow}) at a strictly recorded concentration (C_{pellet}) provides a low-cost MW proxy calibrated directly against HPSEC-MALLS on the same passion-fruit pectin material, avoiding reliance on universal Mark-Houwink constants.

5.3. Proxy Calibration & Acceptance Criteria

To ensure the proxy models generalize, the calibration samples must span the entire operating space. We execute 5 dedicated reactor runs, taking 6 time-resolved aliquots per run (Total $N = 30$).

Run	T (°C)	pH	$d_{\{50\}}$ (μm)	Targeted Domain	Expected State
C1	50	3.0	600	Absolute Min Severity	High M_w , High DE
C2	95	1.5	150	Absolute Max (Stress)	Low M_w , Low DE
C3	75	2.2	300	Interior Trajectory	Moderate M_w , Mod DE
C4	95	3.0	300	High T / Low Acid	Moderate M_w , High DE
C5	50	1.5	150	Low T / High Acid	High M_w , Low DE

Sampling Schedule: Withdrawn at $t = 10, 25, 45, 70, 95, 120$ minutes.

Because time points within a single run are autocorrelated, standard cross-validation will overfit. We mandate **Leave-One-Run-Out (LORO) Cross-Validation**. **Acceptance Metrics:**

1. **RMSEP:** Must be $< 5\%$ absolute for DE/GalA , and < 50 kDa for M_w .
2. **Stratified RMSEP:** The proxy must not fail catastrophically at extremes (e.g., reported separately for low/med/high ranges).
3. **MAE & Bias:** Residuals must be centered at zero.
4. **Uncertainty Coverage:** The proxy's 95% CI must contain the true reference value 95% of the time on unseen runs ($P(y_{\{\text{true}\}} \in CI_{\{95\%, \text{proxy}\}})$).

5.4. Bayesian Integration

The measurement hierarchy resolves into a clean Directed Acyclic Graph (DAG):

$$\text{Kinetic ODE} \rightarrow y_{\{\text{true}\}} \rightarrow y_{\{\text{prepared}\}} \rightarrow z_{\{\text{observed}\}}$$

Where $y = (Y, DE, M_w, X_{\{\text{GalA}\}})$ and $z = (\text{FTIR}, t_{\{\text{flow}\}}, C_{\{\text{ppt}\}}, C_{\{\text{pellet}\}})$. Then the posterior over kinetic parameters is obtained by marginalizing the latent physicochemical state:

$$p(\theta \text{ mid } z, x) \propto p(\theta) \int p(z \text{ mid } y) p(y \text{ mid } \theta, x) dy$$

6. Physical Implementation & Minimum Viable Laboratory (MVP)

The hierarchical measurement architecture directly enables a decentralized, low-capital laboratory strategy. By mathematically separating high-frequency proxies from low-frequency ground-truth anchors, the physical execution requires minimal in-house capital expenditure.

Given the goal is a cheap first-generation reactor, the design explicitly avoids over-instrumentation. The first version maximizes information gained per dollar and per unit of experimental complexity.

6.1. The Lean V1 Sensor Stack

The physical prototype is built around a simple but high-frequency (e.g., 1 Hz) continuous data logging strategy, paired with cheap aliquot-level channels.

Continuous Online Channels:

$$T(t), \text{ pH } pH(t), \text{ RPM } RPM(t), \text{ I } I(t), \text{ m } m(t)$$

Offline Aliquot Channels:

$$NTU(t_{\{\text{sample}\}}), \text{ C } C_{\{\text{ppt}\}}, \text{ FTIR }, t_{\{\text{flow}\}}$$

6.1.1. Continuous Sensor Priority Justification

Priority	Measurement	Cost/Diff.	Why
1	Temperature $T(t)$	Very low	Essential kinetic input.
2	pH $pH(t)$	Low	Essential kinetic input.
3	Agitator RPM	Very low	Makes mixing and load measurements interpretable.
4	Motor Current $I(t)$	Very low	Cheap online proxy for changing slurry load and suspension behavior.

			ior without needing a dedicated torque transducer.
5	Reactor Mass $m(t)$	Low-Mod	Excellent QC for evaporation, sampling mass balance, and S:L ratio drift. A sleeper feature for digital twins.
6	Conductivity $\kappa(t)$	Low	Useful supporting signal (describes liquid environment) but deferred as a second-wave addition.
7	Torque Sensor	Moderate	Better than motor current, but deferred until $I(t)$ proves insufficient.
8	Inline Optics / Density	High	Explicitly deferred until data shows clear necessity.

6.1.2. The Empirical Sensor Fusion Philosophy

The experimental design intentionally pairs simple, ambiguous physical sensors with advanced mathematical calibration. For example:

- **Motor Current ($I(t)$):** At fixed RPM, temperature, and reactor geometry, changes in motor load provide cheap continuous information about the slurry’s evolving physical properties without assuming it directly measures viscosity.
- **Turbidity (NTU):** A cheap optical turbidity measurement on aliquots provides a dense process-state signal. Its interpretation ($NTU = f(\text{particles}, d_{\{50\}}, C_{\{\text{solids}\}})$) is not unique, but together with process history, it powerfully improves the prediction of the latent extraction state.

We record everything at high frequency and ask the data-driven question:

Which cheap signals actually explain the expensive measurements?

If, after 30-50 samples, we discover $M_w \approx f(I, RPM, T, pH, NTU)$, we have successfully discovered a soft sensor empirically. This is exactly the kind of cheap sensor fusion the overarching methodology is designed to exploit.

6.2. Equipment Sourcing Strategy

The project relies on a hybrid execution model, partnering with university core facilities for spectral and chromatographic capabilities:

1. In-House (Own / Operate):

- 1.0 – 2.0 L jacketed glass reactor, mounted on a **load cell**.

- Continuous logging for RTD temperature, pH, agitator RPM, and motor current.
 - Standardized precipitation equipment (centrifuge, vacuum oven).
 - Calibrated analytical balance (Critical: drives the entire mass-basis of the kinetic ODE).
 - Optical turbidity meter (*NTU*).
2. **University / Core Facility Access:**
- Capillary viscometer and UV-Vis spectrophotometer.
 - ATR-FTIR Spectrometer (High-frequency chemical proxy).
 - HPSEC-MALLS (Low-frequency M_w anchor).

Expensive dedicated torque transducers, automated bioreactors, inline FTIR sensors, and sophisticated rheometers are explicitly deferred.

6.3. Analytical Discipline & Traceability

The Bayesian inference engine requires an unbroken relational chain linking the final measurement back to the thermodynamic history of the extraction.

The primary key for all data ingestion is the composite **SampleID** (e.g., B002-R07-S03). This uniquely binds a physical pellet to:

$$t_{\{\text{aliquot}\}} + T(t) + pH(t) + m_{\{\text{liquor}\}} + m_{\{\text{pellet}\}}$$

If this chain is broken, the sample may no longer be suitable for kinetic inference because its process history cannot be assigned reliably.

7. Conclusion & Execution Map

The mechanistic model has been mathematically verified, structural identifiability resolved via D-optimal design, and the measurement architecture strictly defined. By explicitly modeling sample preparation and measurement uncertainty as part of the Bayesian observation hierarchy, the project is insulated from the capital costs and physical biases that typically derail bioreactor optimization.

The reactor is specified here not merely as a machine producing pectin, but as an experimental measurement instrument built specifically for identifying a dynamic model.

The immediate physical step is the **Phase 1 Proxy Calibration**, generating 30 paired observations across 5 OED extreme trajectories to construct the $p(y \mid z)$ measurement likelihoods. Once the chemometric and rheological proxies are validated against the SEC/Titration anchors, the system proceeds to the **10-run OED-selected calibration campaign plus 4 independent validation runs** for final parameter closure.

8. Appendix A: Primer on Bayesian Inference

This appendix provides the conceptual and practical foundation for how Bayesian inference is applied in this digital twin, contrasting it with standard deterministic calibration (e.g., Least Squares).

8.1. A.1 The Conceptual Shift: From Point Estimates to Probability

In traditional process modeling, parameter estimation (calibration) searches for a single “best fit” set of parameters ($\hat{\theta}$) that minimizes the error between model predictions and experimental data.

Bayesian inference asks a fundamentally different question. Instead of searching for one perfect parameter set, it treats parameters as random variables and calculates the **probability of every possible parameter combination** given the observed data.

This is governed by Bayes’ Theorem:

$$p(\theta \text{ mid } D) = \frac{p(D \text{ mid } \theta)p(\theta)}{p(D)}$$

Where:

- $p(\theta)$ is the **Prior**: What we scientifically know about the parameters before running the experiment (e.g., thermodynamic bounds on activation energies).
- $p(D \text{ mid } \theta)$ is the **Likelihood**: The probability of observing our experimental data D if the parameters θ were true. This is generated by running the ODE model.
- $p(\theta \text{ mid } D)$ is the **Posterior**: The updated belief about the parameters after observing the data.

8.2. A.2 Practical Application in the Pectin Digital Twin

In this project, the Bayesian framework solves two critical physical bottlenecks.

8.2.1. 1. Quantifying Uncertainty for Process Optimization

By yielding a full probability distribution ($p(\theta \text{ mid } D)$) rather than a point estimate, the Bayesian model allows us to simulate the reactor thousands of times, each time sampling a different likely parameter set. This produces a predictive distribution (a confidence band) around the expected yield and degradation.

When we later apply optimization algorithms, we do not simply maximize the expected yield. We maximize the **probability of achieving acceptable yield while guaranteeing degradation remains below a critical threshold** (Optimization Under Uncertainty).

8.2.2. 2. The Hierarchical Measurement Model

Standard likelihood functions assume direct, perfect observations of the state variables with simple Gaussian noise. As detailed in the main text, measuring true pectin Molecular Weight (M_w) directly in raw reactor liquor is impossible.

The Bayesian approach naturally handles this through **hierarchical marginalization**. We treat the true physical states (y) as unobserved latent variables, and only condition our inference on the cheap proxy measurements (z):

$$p(\theta \text{ mid } z, x) \propto p(\theta) \int p(z \text{ mid } y)p(y \text{ mid } \theta, x)dy$$

- $p(y \text{ mid } \theta, x)$ is the deterministic ODE.
- $p(z \text{ mid } y)$ is the empirical error distribution of our FTIR and Viscometer proxies, learned explicitly during the Phase 1 Proxy Calibration experiment.

By mathematically integrating over the latent state y , the inference engine correctly absorbs the noise and bias of the cheap proxies, propagating that uncertainty all the way back to the fundamental kinetic parameters.

8.3. A.3 Computational Execution

Because the integral in the denominator of Bayes’ theorem is intractable for a non-linear ODE system, the posterior distribution cannot be calculated analytically.

Instead, we use **Markov Chain Monte Carlo (MCMC)** algorithms (specifically Hamiltonian Monte Carlo / NUTS). These algorithms act as intelligent “walkers” that explore the 9-dimensional parameter space, spending more time in regions where the Prior and Likelihood are high. The resulting trail of samples provides a direct numerical representation of the posterior distribution, entirely bypassing the need for analytical integration.